

Insurance analytics

Neural networks

Katrien Antonio

LRisk - KU Leuven and ASE - University of Amsterdam

May 27, 2019

Acknowledgement

- ▶ Some of the figures in this presentation are taken from
 - Michael A. Nielsen (2015). *Neural networks and deep learning*. Determination Press.
 - all the beautiful work by prof. Taylor Arnold, in particular Chapter 8 of the book *A Computational Approach to Statistical Learning*.

Today's mission

► Today's mission:

- de-mystify neural networks
- sketch different types of neural networks and their applications
- a discussion of [specific considerations](#) to keep in mind when using these predictive modeling techniques with [frequency/severity data](#).

A famous challenge

Recognizing handwritten digits

- ▶ Consider this

504192

- ▶ Humans have a primary visual cortex (V_1) with 140M neurons, with 10s billions connections.
- ▶ Next to V_1 , also V_2 , V_3 , V_4 and V_5 , doing complex image processing.
- ▶ A supercomputer in our head!

A famous challenge

Recognizing handwritten digits




The goal: infer rules for recognizing handwritten digits.

The famous **MNIST** (Modified National Institute of Standards and Technology) database with handwritten digits commonly used for training various image processing systems.

A famous challenge

Recognizing handwritten digits



The winner: well-designed neural nets classify 9 979 out of 10 000 images correctly (in 2013).

Types of neural networks

What's in a name?

- ▶ Different types of neural networks and their applications:
 - **ANN**: Artificial Neural Network
for regression and classification problems, with vectors as input data
 - **CNN**: Convolutional Neural Network
for image processing, image/face/... recognition, with images as input data
 - **RNN**: Recurrent Neural Network
for sequential data such as text or time series.

Types of neural networks

A bit of history

- ▶ 1943: McCulloch and Pitts (1943) with a computational model for neural networks based on maths and algorithms. Paved the way for neural network research in artificial intelligence.
- ▶ 1958: perceptron algorithm for pattern recognition by Rosenblatt.
- ▶ 1969: discovery of two fundamental problems with the computational machines that processed neural networks (e.g. computer power) by Minsky and Papert.
- ▶ 1986: backpropagation algorithm published in Nature.
- ▶ 1990 - 2010: not much happening.
- ▶ 2010 - ... : **AI boom**, driven by (deep) neural networks. Many factors play a role: power of CNNs and RNNs recognized, more computer power, higher demand due to more and more complex data.


Artificial neural networks

A simple neural network

- ▶ De-mystify artificial neural networks (ANNs):
 - a collection of inter-woven linear models
 - extending linear approaches to detect **non-linear** interactions in **high-dimensional** data.

Artificial neural networks

A simple neural network



The goal: predict a scalar response y from scalar input x .

Artificial neural networks

A simple neural network

► Some terminology:

- x is the input layer
- v is output layer
- middle layers are hidden layers
- four neurons: x , z_1 , z_2 and v .

Artificial neural networks

A simple neural network

- ▶ First, apply two independent linear models:

$$z_1 = b_1 + x \cdot w_1$$

$$z_2 = b_2 + x \cdot w_2,$$

using four parameters: two intercepts and two slopes.

- ▶ Next, construct another linear model with the z_j as inputs:

$$\hat{y} := v = b_3 + z_1 \cdot u_1 + z_2 \cdot u_2.$$

Artificial neural networks

A simple neural network


- ▶ Putting it all together:

$$\begin{aligned}v &= b_3 + z_1 \cdot u_1 + z_2 \cdot u_2 \\&= b_3 + (b_1 + x \cdot w_1) \cdot u_1 + (b_2 + x \cdot w_2) \cdot u_2 \\&= (b_3 + u_1 \cdot b_1 + u_2 \cdot b_2) + (w_1 \cdot u_1 + w_2 \cdot u_2) \cdot x \\&= (\text{intercept}) + (\text{slope}) \cdot x.\end{aligned}$$

- ▶ Model is over-parametrized, with infinitely many ways to describe same model.
- ▶ Essentially, still a linear model!

Artificial neural networks

A simple neural network



Artificial neural networks

Activation function

- Capture **non-linear** relationships between x and v by replacing

$$v = b_3 + z_1 \cdot u_1 + z_2 \cdot u_2.$$

with

$$\begin{aligned} v &= b_3 + \sigma(z_1) \cdot u_1 + \sigma(z_2) \cdot u_2 \\ &= b_3 + \sigma(b_1 + x \cdot w_1) \cdot u_1 + \sigma(b_2 + x \cdot w_2) \cdot u_2, \end{aligned}$$

where $\sigma(\cdot)$ is an **activation function**, a mapping from \mathbb{R} to \mathbb{R} .

Artificial neural networks

Activation function

- ▶ Capture **non-linear** relationships between x and v :

$$v = b_3 + \sigma(z_1) \cdot u_1 + \sigma(z_2) \cdot u_2,$$

where $\sigma(\cdot)$ is an **activation function**, a mapping from \mathbb{R} to \mathbb{R} .


- ▶ For example, the rectified linear unit (ReLU) activation function:

$$\text{ReLU}(x) = \begin{cases} x, & \text{if } x \geq 0 \\ 0, & \text{otherwise.} \end{cases}$$

Adding an activation function greatly increases the set of possible relations between x and v !

Artificial neural networks


Activation function



More activation functions: sigmoid, hyperbolic tan, leaky rectified linear unit, maxout.

Artificial neural networks

Artificial vs biological




Artificial neural networks

The architecture

- ▶ Artificial Neural networks (NNs):
 - a collection of neurons
 - organized into an ordered set of layers
 - directed connections pass signals between neurons in adjacent layers
 - **to train**: update parameters describing the connections by minimizing loss function over training data
 - **to predict**: pass \mathbf{x}_i to first layer, output of final layer is \hat{y}_i .
- ▶ The network is *dense* if each neuron in a layer receives an input from all the neurons present in the previous layer; *densely connected*.

Artificial neural networks

The architecture



A basic neural network. Source : <http://blog.christianperone.com>

This is a **feedforward** neural network - no loops!

Artificial neural networks

Terminology

- ▶ Using the NN language:
 - intercept called *the bias*
 - slopes called *weights*
 - L layers in total, with input layer denoted as layer 0
 - use a (from *activation*) to denote the output of a given layer.
- ▶ Let's start with a single layer network (called an artificial neuron).

Artificial neural networks

Terminology





Figure 1: source: andrewjamesturner.co.uk

Artificial neural networks

Terminology

- Use w_{jk}^l to denote the weight for the connection:
 - from neuron k in layer $(l - 1)$
 - to neuron j in layer l .




Artificial neural networks

Terminology

► Use

- b_j^l for the bias of neuron j in layer l
- a_j^l for the activation of neuron j in layer l .



Artificial neural networks

Terminology

- d -dimensional inputs

$$a^0 = x,$$

- weight $w_{j,k}^1$ on k th neuron in layer 0 within j th neuron in layer 1

$$w_j^1 = (w_{j,1}^1, \dots, w_{j,d}^1)$$

- with bias term b_j^1

$$\begin{aligned} z_j^1 &= w_{j,1}^1 \cdot x_1 + \dots + w_{j,d}^1 \cdot x_d + b_j^1 \\ &= \langle w_j^1, x \rangle + b_j^1, \end{aligned}$$

- apply activation function

$$a_j^1 = \sigma(z_j^1),$$

the output of neuron j in layer 1.

Artificial neural networks

Terminology

- ▶ With L layers in total: (in vectorized form)

$$\begin{aligned}z^l &= \langle w^l, a^{l-1} \rangle + b^l \text{ (the weighted input)} \\a^l &= \sigma(z^l),\end{aligned}$$

for l from 1 up to (and including) L .

- ▶ Finally,

$$a^L = \hat{y} = \tilde{\sigma}(z^L),$$

with examples of output activations: identity (regression), sigmoid (binary classification), softmax (multi-class output).

Stochastic gradient descent

Basic idea - GD

- ▶ We want to find

$$\min_w f(w),$$

then **gradient descent** updates as follows (cfr. lecture on *boosting methods*)

$$w_{\text{new}} = w_{\text{old}} - \eta \cdot \nabla_w f(w_{\text{old}}),$$

with **learning rate** η . Move in the direction the function locally decreases the fastest!

- ▶ A good choice for NNs because faster second-order methods involve the Hessian and this is infeasible when having tons of parameters.

Stochastic gradient descent

From GD to SGD

- ▶ Let $\mathcal{L}(w; y_i, x_i)$ be the loss function and w the parameters. For example,

$$\begin{aligned}\mathcal{L}(w; y_i, x_i) &= \frac{1}{2n} \sum_i (\hat{y}_i(w) - y_i)^2 \\ &= \frac{1}{n} \sum_i \mathcal{L}_i(w; y_i, x_i).\end{aligned}$$

- ▶ With the gradient descent update rule

$$w_{\text{new}} = w_{\text{old}} - \eta \cdot \nabla_w \mathcal{L}(w),$$

where $\nabla_w \mathcal{L}(w) = \frac{1}{n} \sum_i \nabla_w \mathcal{L}_i$.

Stochastic gradient descent

Mini-batches

- ▶ Stochastic gradient descent picks randomly m training inputs x_1, \dots, x_m , a **mini-batch**:

$$\frac{\sum_{j=1}^m \nabla_w \mathcal{L}_j}{m} \approx \frac{\sum_i \nabla_w \mathcal{L}_i}{n} = \nabla_w \mathcal{L}.$$

- ▶ Thus,

$$\nabla_w \mathcal{L} \approx \frac{1}{m} \sum_{j=1}^m \nabla \mathcal{L}_j(w).$$

- ▶ Estimate the gradient only on a small subset of the entire training set.

Stochastic gradient descent

- ▶ Partition input randomly into disjoint groups $M_1, M_2, \dots, M_{n/m}$.
- ▶ Updates:

$$\begin{aligned}w_{k+1} &= w_k - \frac{\eta}{m} \sum_{i \in M_1} \nabla \mathcal{L}_i \\w_{k+2} &= w_{k+1} - \frac{\eta}{m} \sum_{i \in M_2} \nabla \mathcal{L}_i \\&\vdots \\w_{k+n/m+1} &= w_{k+n/m} - \frac{\eta}{m} \sum_{i \in M_{n/m}} \nabla \mathcal{L}_i.\end{aligned}$$

Speed up the process of doing gradient descent. Going through the entire data set is called an **epoch**.

Backward propagation of errors

- ▶ Compute the gradient of the loss function wrt all trainable parameters:
 - tons of parameters
 - need for efficient algorithm to calculate gradient
 - generic algorithm usable for arbitrary number of layers and neurons in each layer.
- ▶ The strategy (Rumelhart et al., 1986, Nature)

backwards propagation of errors or **backpropagation**

uses chain rule for derivatives.

Backward propagation of errors

- ▶ With a loss function \mathcal{L} , e.g. squared error loss

$$\mathcal{L}(y, a^L) = \frac{1}{2n}(y - a^L)^2.$$

- ▶ Equations defining backpropagation: (recall: $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$)

starting point - gradient terms of b_j^l

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial b_j^l} &= \frac{\partial \mathcal{L}}{\partial z_j^l} \cdot \frac{\partial z_j^l}{\partial b_j^l} \\ &= \frac{\partial \mathcal{L}}{\partial z_j^l} \cdot 1 = \frac{\partial \mathcal{L}}{\partial z_j^l}.\end{aligned}$$

Thus, problem centers around derivatives wrt z^l !

Backward propagation of errors

- ▶ With a loss function \mathcal{L} , e.g. squared error loss

$$\mathcal{L}(y, a^L) = \frac{1}{2n}(y - a^L)^2.$$

- ▶ Equations defining backpropagation: (recall: $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$)

starting point - gradient terms of weights w_{jk}^l

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_{jk}^l} &= \frac{\partial \mathcal{L}}{\partial z_j^l} \cdot \frac{\partial z_j^l}{\partial w_{jk}^l} \\ &= \frac{\partial \mathcal{L}}{\partial z_j^l} \cdot a_k^{l-1}.\end{aligned}$$

Thus, problem centers around derivatives wrt z^l !

Backward propagation of errors

- The derivatives in layer L are straightforward to calculate (with $a_j^L = \sigma(z_j^L)$)

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial z_j^L} &= \sum_k \frac{\partial \mathcal{L}}{\partial a_k^L} \cdot \frac{\partial a_k^L}{\partial z_j^L} \\ &= \frac{\partial \mathcal{L}}{\partial a_j^L} \cdot \frac{\partial a_j^L}{\partial z_j^L} \\ &= \frac{\partial \mathcal{L}}{\partial a_j^L} \cdot \sigma'(z_j^L).\end{aligned}$$

An equation for the error in the output layer!

Backward propagation of errors

- ▶ The derivatives wrt z^l can be written as a function of derivatives wrt z^{l+1} .
- ▶ Using the chain rule

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial z_j^l} &= \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{l+1}} \cdot \frac{\partial z_k^{l+1}}{\partial z_j^l} \\ &= \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{l+1}} \cdot w_{kj}^{l+1} \sigma'(z_j^l),\end{aligned}$$

where we use

$$z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} \sigma(z_j^l) + b_k^{l+1}.$$

Backward propagation of errors

► The vanilla implementation of the backpropagation algorithm:

1. **Input** x , set a^0 for the input layer.

2. **Feedforward**: for each $l = 2, 3, \dots, L$ compute

$$z^l = \langle w^l, a^{l-1} \rangle + b^l \text{ and } a^l = \sigma(z^l).$$

3. **Output error**: compute $\frac{\partial \mathcal{L}}{\partial z_j^L} = \frac{\partial \mathcal{L}}{\partial a_j^L} \cdot \sigma'(z_j^L)$.

4. **Backpropagate the error**: for each $l = L - 1, \dots, 2$ compute $\frac{\partial \mathcal{L}}{\partial z_j^l} = \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{l+1}} \cdot w_{kj}^{l+1} \sigma'(z_j^l)$.

5. **Output**: the gradient of the loss function $\frac{\partial \mathcal{L}}{\partial w_{jk}^l} = \frac{\partial \mathcal{L}}{\partial z_j^l} \cdot a_k^{l-1}$ and $\frac{\partial \mathcal{L}}{\partial b_j^l} = \frac{\partial \mathcal{L}}{\partial z_j^l}$.

Improving SGD

- ▶ Several improvements covered in the literature
 - to reduce **overfitting**
 - to reduce getting stuck in **local saddle points**
 - to **converge** in a smaller number of epochs.

Improving SGD

Reduce overfitting

► Methods to reduce overfitting:

- different weight initialization
- early stopping
- regularization
- dropout.

Improving SGD

Reduce overfitting: weight initialization, early stopping

- ▶ Update the **weight initialization** \rightsquigarrow might drastically improve performance of a model with many layers.
- ▶ **Early stopping** to prevent overfitting:
 - calculate validation performance after each epoch
 - stop when this no longer improves
 - NNs are motivated by an optimization problem, but do not attempt to solve the optimization task.

Improving SGD

Reduce overfitting: regularization

- ▶ Prevent overfitting via **regularization** (cfr. lecture on Lasso and friends)
 - add (e.g.) ℓ_2 -norm

$$f_\lambda(w, b) = f(w, b) + \frac{\lambda}{2} \cdot \|w\|_2^2,$$

with weights w and bias terms b , where only weights are penalized

- with gradient

$$\nabla_w f_\lambda = \nabla_w f(w, b) + \lambda \cdot w$$

such that

$$\begin{aligned} w_{\text{new}} &\leftarrow w_{\text{old}} - \eta \cdot \nabla_w f_\lambda(w_{\text{old}}) \\ &\leftarrow w_{\text{old}} - \eta \cdot [\nabla_w f(w_{\text{old}}) + \lambda \cdot w] \\ &\leftarrow [1 - \eta \cdot \lambda] \cdot w_{\text{old}} - \eta \cdot f(w_{\text{old}}). \end{aligned}$$

Improving SGD


Reduce overfitting: dropout

► Dropout:


- randomly set activations to zero, with fixed probability p
- both in forward propagation as well as backpropagation
- only in training, all nodes turned on during prediction.

Improving SGD and regularization

Dropout



(a) Standard Neural Net



(b) After applying dropout.

Figure 5: Dropout - source: <http://blog.christianperone.com/>

Improving SGD

Saddle points & faster convergence

- ▶ There is a fast proliferation of improved optimization algorithms.
- ▶ Examples:
 - momentum
 - RMSprop
 - adagrad, adamdelta
 - adam, adamax
 - natural gradient descent

Where to finetune your ANN?

- ▶ A list of [tuning parameters/architecture choices](#):
 - learning rate η
 - regularization parameter λ
 - number of epochs
 - mini-batch size
 - number of layers
 - number of hidden neurons per layer
 - activation functions, choice of optimization strategy.


Classification with Artificial neural networks

Back to the famous challenge

- ▶ **The goal:** try to recognize individual hand-written digits
 - take e.g. 28 by 28 greyscale image
 - $784 = 28 \times 28$ input neurons, with intensities between 0 (white) and 1 (black)
 - output layer has 10 layers
- neuron with highest activation value fires.

Classification with artificial neural networks

Back to the famous challenge



Classification with artificial neural networks

- ▶ Use a multi-valued output layer for classification tasks.
- ▶ One-hot encoding applied to output vector y

$$\begin{pmatrix} 2 \\ 4 \\ \vdots \\ 1 \end{pmatrix} \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}.$$

- ▶ Use **softmax** function as activation in final layer

$$\begin{aligned} a_j^L &= \text{softmax}(z_j^L) \\ &= \frac{e^{z_j}}{\sum_k e^{z_k}}, \end{aligned}$$

for the j th output neuron. The last layer is then easily interpreted as sequence of probabilities.

Classification with artificial neural networks

- ▶ Instead of using squared error loss, use categorical cross-entropy

$$\mathcal{L}(a^L, y) = -\frac{1}{n} \sum_i \sum_k y_{ik} \cdot \log(a_k^L),$$

as the loss function, where i runs over training data and k runs over the class labels.

- ▶ Think: softmax output layer and negative log-likelihood cost function.
- ▶ Work out the derivatives to set-up the backpropagation with this loss function.

Convolutional neural networks

The motivation

- ▶ Prediction tasks with images as inputs are a popular application of NNs.
- ▶ With a limited number of input pixels, put a weight on each pixel and use ANN.
- ▶ With large images use convolutional layers.

Convolutional neural networks

A convolution?

- ▶ The discrete convolution between f and g is

$$(f * g)(x) = \sum_t f(t) \cdot g(x + t).$$

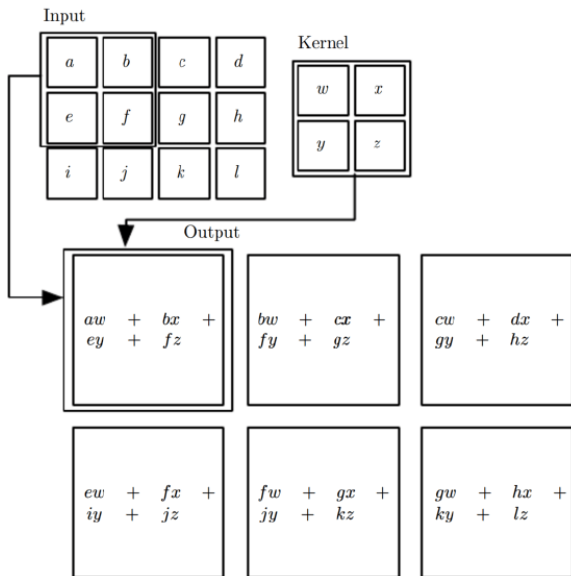
- ▶ With 2-dimensional signals (e.g. images), consider 2D-convolutions

$$(K * I)(i, j) = \sum_{m, n} K(m, n) \cdot I(i + m, j + n),$$

with K a convolutional kernel applied to a 2D signal (or image) I .

Convolutional neural networks

A convolution pictured



Convolutional neural networks

- ▶ With large images use convolutional layers:
 - apply small set of weights to subsections of the image
 - same weights across the image
 - use a kernel matrix, e.g. for a black and white image

$$K = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}.$$

Take pixel value and subtract from this the pixel value to its immediate lower right.

Convolutional neural networks

- ▶ In a CNN:
 - apply several such kernels in a layer
 - with weights learned during training of the algorithm
 - different convolutions pick up different features (e.g. edges, texture, basic object types).
- ▶ Pooling layer: simplify information in output from convolutional layer.
- ▶ With input image in color: kernel matrix K is 3-dimensional array with weights applied to each color channel.

Recurrent neural networks

- ▶ To infer sequential data such as text or time series.
- ▶ Mathematically,

$$\begin{aligned}h_t &= Wx_t + b + Uh_{t-1} \\&= Wx_t + b + UWx_{t-1} + Ub + U^2h_{t-2} \\&= \dots\end{aligned}$$

Recurrent neural networks

Unrolled




Figure 19: Unrolled representation of a RNN. Source : Understanding LSTM Networks by Christopher Olah - <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Yes, we CANN!



- ▶ Starting from the Poisson regression model

$$N_i \sim \text{POI}(\lambda(\mathbf{x}_i, d_i)),$$

with d_i exposure and \mathbf{x}_i the features for insured i .

- ▶ Build e.g.
 - a fully-connected single hidden layer feed-forward neural network with 20 hidden neurons for modelling the regression function
 - a more general architecture with K hidden layers.

Yes, we CANN!



A shallow (left) and a deep (right) neural network.

Yes, we CANN!

► Some comments:

- output layer has a single neuron, and exponential activation function
- some preprocessing steps:

ordered categorical features as continuous

dummy coding or one-hot encoding to construct binary representation of nominal features

put continuous covariates on the same scale

- use Poisson loss function in keras.

Yes, we CANN!

Listing 3: R script for fitting networks in Keras

```
1 library(keras)
2
3 model <- keras_model_sequential()
4 model %>%
5   layer_dense(units = q1, activation = 'tanh', input_shape = c(ncol(Xlearn))) %>%
6   layer_dense(units = 1, activation = k_exp)
7
8 summary(model)
9
10 model %>% compile(
11   loss = 'poisson',
12   optimizer = 'sgd'
13 )
14
15 fit <- model %>% fit(Xlearn, learn$ClaimNb, epochs=100, batch_size=10000)
```

keras implementation of a neural network for frequencies.

Yes, we CANN!

- ▶ Explore the following links
 - Case study: French MTPL claims
 - Insights from inside neural networks
 - Data analytics for non-life insurance pricing
 - Editorial: Yes we CANN!